

When Better QAOA Schedules Depend on the Simulator: Truncation-Induced Rank Reversals on a 55-Qubit QOBLIB Benchmark

Anonymous Authors

Research draft; author names and affiliations withheld

Abstract

Approximate tensor-network simulators make variational quantum algorithms accessible beyond the statevector limit, but algorithm rankings are often reported at a single approximation setting. We test whether such a ranking is stable for depth-15 Quantum Approximate Optimization Algorithm (QAOA) schedules on a real Maximum Independent Set benchmark. A four-parameter nonlinear ramp was selected on full QOBLIB ES60FST instances and deployed without retuning to `es60fst02`, a 186-vertex graph reduced by a published pipeline to 55 qubits and 91 edges. At the released Aer matrix product state (MPS) setting (bond dimension 64, truncation threshold 10^{-3}), 15 paired jobs and 15,000 shots per method yielded 101 best-known solution (BKS) hits for the nonlinear ramp and 41 for the published linear ramp: probabilities 0.00673 and 0.00273, a paired difference of 0.00400 (bootstrap 95% CI [0.00247, 0.00553], exact sign-flip $p = 2.44 \times 10^{-4}$). However, tightening only the truncation threshold to 10^{-4} reversed the BKS ranking: in 10,000 shots, the linear and nonlinear ramps reached 0.0153 and 0.0109, respectively. An intermediate threshold 3×10^{-4} still favored the nonlinear ramp, while raising the bond cap from 64 to 96 at 10^{-3} did not reverse the ranking. Feasible probability and probability of BKS-1-or-better outcomes remained higher for the nonlinear ramp across all tested settings. Thus, a rare-event optimum-hit metric can rank transferred QAOA schedules differently as MPS truncation is tightened even when broader distributional metrics agree. Exact-state calibration on the 12- and 15-qubit donor kernels confirms that threshold 10^{-3} produces total-variation error 0.05–0.07 and schedule-dependent positive BKS bias; at 10^{-6} all state fidelities exceed 0.99983 and absolute BKS error is below 2.3×10^{-4} . An independent cuTensorNet implementation reproduces exact small-kernel states at unit fidelity. On the 55-qubit circuit, its spectral-order bond-128 MPS changes from 15 versus 6 nonlinear/LR BKS hits at cutoff 10^{-3} to 11 versus 12 at 10^{-4} in 5,000 shots per method. Classical controls dominate both: HiGHS proves BKS 88 in 0.036 seconds, and randomized minimum-degree greedy reaches BKS in 7.55% of 15,000 full-graph starts. The result motivates convergence sweeps and classical controls as required components of approximate-simulator optimization benchmarks.

1 Introduction

QAOA alternates problem and mixer Hamiltonians with variational angles chosen by a classical optimization loop [1]. Parameter transfer amortizes that loop by learning angles on donor instances and deploying them on larger acceptors [2, 3]. This strategy is especially attractive when the target circuit is too large for exact statevector simulation. In practice, such targets are often evaluated with approximate tensor networks, whose runtime and accuracy are controlled by bond dimensions and truncation thresholds [10, 12].

The optimization literature usually asks which QAOA schedule is better. The simulation literature asks how runtime and aggregate error scale with an approximation parameter. We ask a question at their intersection: *does the simulator preserve the ordering of QAOA schedules on the observable used to declare a winner?* This is not guaranteed. The probability of an optimum may live in a small tail of the output distribution; a small state approximation error can therefore cause a large relative error in that rare-event probability without visibly changing average energy or feasibility.

We study this issue on the Quantum Optimization Benchmarking Library (QOBLIB), a curated, submission-oriented framework designed for reproducible comparison of quantum and classical optimization methods [8]. We reproduce the released depth-15 Aer/MPS pipeline for the `es60fst02` Maximum Independent Set (MIS) instance, including graph reduction, Hamiltonian construction, sample unfolding, and a validity-only filter. We compare its published linear-ramp schedule with a low-dimensional nonlinear ramp selected on other full ES60FST instances.

Our contributions are:

1. a train-validation-test protocol over four complete QOBLIB graphs, ending in a 55-qubit, depth-15 blind benchmark with no repair or target retuning;
2. an equal-budget comparison of evolutionary and random schedule search, showing that the random-search champion, not the evolutionary operator, transfers best;
3. 15 paired MPS jobs at the published simulator setting, with bootstrap intervals and an exact paired randomization test; and
4. a bond/threshold sensitivity audit that identifies a truncation-induced BKS rank reversal while near-optimal and feasible-mass rankings remain stable.

We do not claim quantum advantage, an exact simulation of the 55-qubit state, or universal superiority of the nonlinear ramp. The main claim is about benchmark validity: schedule rankings based on rare optimum hits can depend on an approximate simulator’s truncation setting.

2 Background and related work

Parameter transfer and schedule reduction. QAOA parameters concentrate across related random graphs, motivating donor selection, representation learning, and direct parameter transfer [2, 3]. Weighted-problem studies show that rescaling phase angles can restore transfer when coefficient distributions change [4]. Linear-ramp QAOA compresses $2p$ angles to two increments and has been studied at larger depths [5]. Recent constrained-QAOA work promotes penalty schedules and piecewise ramps to variational parameters [6]. Our nonlinear ramp is intentionally modest; it is a vehicle for testing ranking stability, not the central theoretical novelty.

Evolutionary variational optimization. Evolutionary methods are natural for nonconvex, noisy quantum objectives [7]. We use a fixed-population genetic search and compare it with uniform random search under exactly the same number of candidate evaluations. The negative result—matched random search finds the stronger transferring schedule—prevents attributing the observed gain to evolution itself.

Approximate tensor-network simulation. MPS simulation compresses a many-qubit state by truncating Schmidt spectra; accuracy depends on circuit ordering, entanglement growth, the maximum bond dimension, and discarded-weight policy [10]. Recent work studies tensor-network surrogates for variational circuits and the compute-accuracy frontier of approximate simulators [11, 12]. Exact tensor-network sampling methods have also reached structured QAOA circuits beyond conventional statevector sizes [13]. These studies motivate careful convergence analysis, but they do not establish that optimization algorithm rankings are invariant under the approximation controls used in a submitted benchmark.

3 Benchmark protocol

3.1 Instances and frozen split

Table 1 lists the four full QOBLIB ES60FST MIS instances. The split was frozen before evaluation of any new schedule on **es60fst02**. Best-known values are exact Gurobi optima in the QOBLIB submission records.

We use the released reduction implementation: sequential exact degree-0/1 rules, exact degree-2 folding, recorded pruning of vertices above degree four, and a final exact cleanup. Pruning is heuristic; all decisions are recorded

Table 1: Full QOBLIB instances and quantum kernels at maximum retained degree four. V/E are original vertices/edges; n_q/m_q are reduced vertices/edges.

Instance	V	E	n_q	m_q	Role / BKS
es60fst01	123	159	15	21	train / 60
es60fst03	113	142	12	16	train / 55
es60fst04	162	238	45	69	validation / 78
es60fst02	186	280	55	91	blind test / 88

so samples can be unfolded. On **es60fst02**, 37 vertices are forced selected, 38 forced excluded, 28 folds occur in four batches, and five vertices are heuristically pruned. The published workflow also uses this 55-qubit kernel and reaches BKS 88 under Aer MPS [9].

3.2 Hamiltonian, circuit, and schedules

For graph $G = (V, E)$, the MIS QUBO is

$$C(x) = - \sum_{i \in V} x_i + \lambda \sum_{(i,j) \in E} x_i x_j, \quad \lambda = 1.5. \quad (1)$$

After $x_i = (1 - Z_i)/2$, QAOA alternates the Ising cost operator and $H_M = \sum_i X_i$. Both methods use depth $p = 15$, maximum-quadratic-coefficient Hamiltonian scaling from the released code, an initial $|+\rangle^{\otimes n}$ state, and identical measurement conventions. The untranspiled test circuit contains 1,365 RZZ, 825 RZ, 825 RX, 55 Hadamard, and 55 measurement operations.

The published linear ramp is

$$\beta_k = 0.7 \frac{p - k + 1}{p}, \quad \gamma_k = 0.4 \frac{k}{p}. \quad (2)$$

The candidate family generalizes the ramp with two powers:

$$\beta_k = \Delta_\beta \left(\frac{p - k + 1}{p} \right)^{a_\beta}, \quad (3)$$

$$\gamma_k = \Delta_\gamma \left(\frac{k}{p} \right)^{a_\gamma}. \quad (4)$$

The four-dimensional genome lies in $\Delta_\beta \in [0.15, 1.20]$, $\Delta_\gamma \in [0.05, 1.00]$, and $a_\beta, a_\gamma \in [0.30, 2.50]$. The frozen random-search champion is (0.64247, 0.75939, 1.77679, 0.99172); the evolutionary champion is (0.51750, 0.77197, 1.07734, 1.75435).

3.3 Search, selection, and metrics

Three independent search replicates use 20 candidates for six generations, or 120 candidate evaluations per method. Each candidate receives 256 shots on each training instance. Random search receives exactly 120 independently drawn candidates. A finite-shot robust score combines Wilson lower bounds on BKS, BKS-1, and feasible probabilities with feasible quality mass; its mean and worst-instance

values are combined 3:1. One validation pass on `es60fst04` selects one champion per search method. Test outcomes are never used for schedule selection.

For every reduced sample, the released decoder reverses folds and forced decisions. We retain the full sample only if it is already an independent set. There is no constraint repair, greedy fill, local search, or archived-solution fallback. We report

- BKS hit probability $\mathbb{P}(|S| \geq 88)$;
- near-BKS probability $\mathbb{P}(|S| \geq 87)$;
- feasible probability; and
- quality mass $\mathbb{E}[\mathbf{1}_{\text{feasible}} \min(|S|/88, 1)]$.

The confirmatory run uses 15 paired simulator seeds and 1,000 shots per method and seed. Paired mean differences receive a 20,000-draw bootstrap interval and an exact two-sided sign-flip test over all 2^{15} paired sign assignments.

3.4 Independent backend and classical controls

We export the frozen native circuits to NVIDIA cuTensorNet 26.6 [14]. Direct contraction reproduces Qiskit statevectors for both schedules on both 12- and 15-qubit kernels with fidelity one to displayed precision. The 55-qubit exact sampler fails during network preparation, so all completed large-circuit results are explicitly approximate MPS results. We use a Fiedler-vector qubit ordering, exact-equivalent before truncation, that lowers edge bandwidth from 47 to 10 and maximum linear cut from 42 to 9. We test bond 128 with discarded-weight cutoffs 10^{-3} and 10^{-4} using 5,000 shots per schedule. Samples are inverse-permuted before the unchanged decoder is applied.

Classical controls operate on the original full graph and on the released 55-variable kernel. An exact binary MILP is solved with SciPy/HiGHS at zero MIP gap. A randomized minimum-residual-degree heuristic repeatedly selects a minimum-degree vertex with random tie breaking, deletes it and its neighbors, and uses no archived solution. Both heuristic variants receive 15,000 starts.

3.5 Simulator settings and audit

The primary setting exactly follows the submission: Qiskit Aer MPS, maximum bond dimension 64, truncation threshold 10^{-3} , and native RZ/RZZ/RX circuits with no hardware-basis transpilation. This detail matters because a gate decomposition changes when MPS truncation is applied. A first transpiled pass was preserved as a protocol audit and excluded from the main result.

We then hold circuits and shots fixed while testing thresholds 10^{-3} , 3×10^{-4} , and 10^{-4} at bond 64, plus bonds 64 and 96 at thresholds 10^{-3} and 10^{-4} . Audit points use 10,000 shots per principal method; the primary $10^{-3}/64$ point aggregates 15,000 shots.

To calibrate the direction and magnitude of approximation error, we also run complete statevectors on the real 12- and 15-qubit training kernels. For both principal

Table 2: Blind `es60fst02` results at bond 64 and threshold 10^{-3} ; 15,000 shots per method.

Method	BKS	BKS-1	Feasible	Quality
Published LR	.00273	.11213	.50127	.48962
Evolutionary	.00207	.06813	.89987	.87202
Matched random	.00673	.15700	.90180	.87865

schedules we compare exact probabilities with MPS states at bond dimensions 16, 32, 64, and 128 and thresholds from 10^{-3} to 10^{-6} . We report state fidelity, total-variation distance, Jensen-Shannon divergence, and signed errors in BKS, near-BKS, feasibility, and quality mass. These are probability-exact comparisons with no measurement noise.

4 Results

4.1 Search method is not the source of the gain

On the 45-qubit validation kernel, the evolutionary champion raises feasible probability from 0.561 to 0.930 but lowers BKS probability from 0.040 to 0.017. The matched-random champion attains 0.941 feasibility, 0.051 BKS, and 0.475 near-BKS. This result was fixed before the corrected blind pass. It shows that the nonlinear family contains transferable schedules, but the evolutionary operator does not outperform a matched random search at this budget.

4.2 Published MPS setting

Table 2 and Fig. 1 summarize 45 method-jobs on the blind instance. All methods find BKS 88 in every aggregate. The random-search nonlinear ramp improves all four reported rates relative to the published linear ramp. It produces 101 BKS hits versus 41, and its BKS advantage is positive in 13 seeds with two ties and no losses. The estimated shot count for a 95% chance of at least one BKS sample, $\lceil \log(0.05) / \log(1 - p_{\text{BKS}}) \rceil$, falls from 1,095 to 444.

For matched random minus published LR, the paired BKS difference is 0.00400 (bootstrap 95% CI $[0.00247, 0.00553]$, sign-flip $p = 0.000244$). Corresponding differences are 0.04487 for near-BKS, 0.40053 for feasibility, and 0.38904 for quality mass; all three have sign-flip $p = 6.10 \times 10^{-5}$. The evolutionary champion improves feasibility and quality mass but not BKS or near-BKS, reinforcing the negative optimizer result.

4.3 Truncation-induced rank reversal

Figure 2 shows the central finding. At bond 64 and threshold 10^{-3} , the nonlinear ramp has the larger BKS rate. At 3×10^{-4} it still leads, 0.0161 to 0.0116. At 10^{-4} the order reverses: published LR reaches 0.0153 and the nonlinear ramp 0.0109. The shot-noise standard error of a rate near




Figure 1: Blind metrics at the released Aer/MPS setting. LR is the published linear ramp; ES and RS are the frozen evolutionary- and random-search nonlinear ramps. Bars aggregate 15 paired jobs and 15,000 shots per method.

0.01 with 10,000 shots is approximately 0.001, smaller than the observed 0.0044 gap.

The same threshold tightening does not reverse the broader metrics. At $10^{-4}/64$, nonlinear versus LR is 0.1621 versus 0.1546 for near-BKS and 0.7351 versus 0.4809 for feasibility. At $10^{-4}/96$ it is 0.1808 versus 0.1710 near-BKS and 0.7121 versus 0.4799 feasible. The optimum-hit tail is thus more ranking-sensitive than near-optimal or feasible mass.

Increasing the bond cap alone does not explain the effect. At threshold 10^{-3} and bond 96, nonlinear still leads BKS 0.0070 to 0.0019, near-BKS 0.1661 to 0.1188, and feasibility 0.8925 to 0.5055. At threshold 10^{-4} and bond 96, the BKS comparison is 0.0134 versus 0.0152, again favoring LR, while near-BKS and feasibility continue to favor nonlinear. The threshold is the dominant control in the tested range.

Runtime makes this audit consequential rather than cosmetic. At bond 96, tightening the threshold from 10^{-3} to 10^{-4} increases the two-circuit batch from about 38 seconds to about 347 seconds on the local CPU. An attempted bond-128, 10^{-6} three-method batch exceeded 20 minutes and was terminated without producing a result. Single-setting reports therefore create a real temptation to use a cheap but unconverged point.

4.4 Exact-state calibration isolates threshold bias

The small-kernel calibration supplies ground truth for the same schedules and reduction family. On **es60fst01**, exact BKS probabilities are 0.68610 for LR and 0.92898 for the nonlinear ramp. At threshold 10^{-3} , MPS shifts them by +0.01577 and +0.03417, respectively. On **es60fst03**, exact BKS probabilities are 0.73583 and 0.95530; corresponding MPS biases are +0.00744 and +0.01849. Thus, the released threshold overstates BKS probability for both schedules but more strongly for the nonlinear ramp.

Figure 3 shows convergence. At 10^{-3} , total-variation distance is 0.0518/0.0509 for LR and 0.0742/0.0702 for nonlinear on the two kernels. At 10^{-6} all four state fidelities exceed 0.99983, total variation is below 0.00266, and ab-

solute BKS error is below 0.00023. For shared thresholds, results are numerically identical at bonds 32, 64, and 128. Bond 16 differs only for the nonlinear **es60fst01** circuit at 10^{-4} , where total variation changes by less than 0.0005. Thus, once the cap reaches 32, the discarded-weight threshold controls the error. The nonmonotonic signed BKS error for LR also shows why convergence of a rare observable cannot be inferred solely from a monotonically improving global state metric.

4.5 Independent backend and classical controls

At spectral ordering and bond 128, independent cuTensorNet MPS yields 15 versus 6 BKS hits for nonlinear versus LR at cutoff 10^{-3} in 5,000 shots per method: rates 0.0030 and 0.0012 (two-sided Fisher exact $p = 0.078$). Tightening to 10^{-4} removes that directional advantage: counts are 11 versus 12, rates 0.0022 and 0.0024 ($p = 1.0$). Near-BKS changes from 0.0820 versus 0.0424 to 0.0782 versus 0.0764, while feasible probability remains higher for nonlinear, 0.6232 versus 0.4840 and 0.6352 versus 0.4838. Thus the independent backend does not reproduce a significant LR advantage at the tighter point, but it does reproduce loss of the nonlinear BKS advantage while the feasible ranking remains stable.

The classical comparison rules out a solver-advantage interpretation. HiGHS proves BKS 88 with zero gap in 0.036 seconds. On the full 186-vertex graph, randomized minimum-degree greedy produces 1,133 BKS solutions in 15,000 starts (0.07553) in 19.95 seconds. On the same QOBLIB kernel used by QAOA, it produces 6,063 (0.4042) in 2.21 seconds. These BKS rates are 11.2 and 60.0 times the best QAOA rate. One third of the measured three-method QAOA batch wall time is about 183 seconds, so the corresponding classical heuristic runs are also about 9.2 and 82.9 times faster. The contribution is therefore diagnostic methodology, not quantum or heuristic optimization advantage.

Schedule ranking is observable-dependent and truncation-sensitive (bond 64)




Figure 2: Observable-dependent convergence at fixed bond dimension 64. Tightening the threshold reverses the BKS ranking between 3×10^{-4} and 10^{-4} , while near-BKS and feasible rankings remain unchanged. The 10^{-3} point uses 15,000 shots; the other points use 10,000 shots per method.

Exact-state calibration on real QOBLIB kernels (bond 128)




Figure 3: Exact-state MPS calibration on the 15-qubit `es60fst01` and 12-qubit `es60fst03` kernels at bond 128. Left: total-variation distance from the exact output distribution. Right: signed error in BKS probability. Tightening the truncation threshold drives both errors toward zero, but the rare-event bias is schedule-dependent and need not be monotone.

5 Discussion

5.1 Why a rank reversal is plausible

MPS compression repeatedly removes small Schmidt components. BKS probability is a rare-event functional of the final distribution rather than a smooth global norm. A schedule can therefore appear superior because the simulator retains or discards a small set of amplitudes differently. Tightening the threshold changes BKS probability by factors of several for both schedules, yet feasibility changes far less for LR and remains much higher for the nonlinear ramp. This pattern is consistent with rare optimum strings being more sensitive than broad feasible regions. The exact-state calibration supports this mechanism directly: global distribution error falls smoothly as the threshold tightens, while signed BKS error is schedule-specific and crosses zero for

LR.

The experiment does not identify which schedule is closer to the exact 55-qubit state. Convergence would require tighter settings, an independent exact sampler, or quantum-hardware evidence. The correct conclusion is not that LR is truly better at 10^{-4} , but that neither BKS ranking is established by a single approximate setting.

The independent backend narrows this statement: changing implementation, qubit ordering, bond cap, and cut-off changes absolute tail probabilities and removes either schedule’s stable BKS advantage. In contrast, the feasible-mass ordering survives. Qubit ordering must therefore be reported alongside the usual MPS truncation controls.

5.2 Implications for optimization benchmarks

We recommend that approximate-simulator submissions report:

1. native circuit, qubit ordering, bond cap, discarded-weight threshold, and whether a basis transpilation was applied;
2. at least three accuracy points bracketing the reported setting;
3. ranking stability for every headline observable, not only average energy or a simulator fidelity proxy;
4. raw counts, shot budgets, and confidence intervals for rare-event metrics; and
5. a distinction between conclusions stable across settings and those conditional on one released benchmark configuration.

These requirements complement QOBLIB’s standardized reporting objective [8]. They also matter for transfer learning: an approximate simulator can influence not only final evaluation but validation-time model selection. Our champions were deliberately frozen after a documented code audit, and the first transpiled run remains archived rather than silently deleted.

5.3 Limitations

The study uses one four-instance graph family and one MIS reduction policy. The high-degree pruning step is heuristic, although all reported samples are valid on the original graph and BKS values refer to the full instances. Only two principal schedules receive the full threshold sweep. For the 55-qubit circuit we do not compute state fidelity, discarded-weight accumulation, or exact probabilities. Direct cuTensorNet contraction matches all small exact controls but its 55-qubit sampler returns an internal preparation error; the independent large-circuit results remain approximate. Both MPS simulators are noiseless and neither constitutes hardware evidence. Exact-state calibration is limited to the 12- and 15-qubit members of the same family and cannot certify convergence of the 55-qubit state. Finally, the 10,000-shot sensitivity points are single simulator jobs, so their binomial precision is high but job-to-job implementation effects are not separately estimated.

6 Reproducibility

The experiment records Python, Qiskit, Aer, cuTensorNet, NumPy, operating-system metadata, the cloned baseline commit, and a SHA-256 digest of the frozen protocol. Training, validation, blind testing, the discarded transpilation audit, every fidelity point, independent-backend job, and classical run are separate JSON artifacts. Checkpoints are written after long jobs. Eight integrity tests verify the published ramp formula, split disjointness, the 186-to-55 reduction, disabled repair, balanced blind artifacts, exact

calibration, classical feasibility, cuTensorNet unit-fidelity controls, and both 5,000-shot independent points. Analysis scripts regenerate CSV tables and all three figures; the exact-calibration artifact contains 56 complete state comparisons.

7 Conclusion

On a real 186-vertex QOBLIB instance executed as a 55-qubit, depth-15 QAOA circuit, a transferred nonlinear ramp appears to improve optimum-hit probability by 2.46 times at the released MPS setting. That headline does not survive a tighter truncation threshold: the optimum-hit ranking reverses, while near-optimal and feasible-mass rankings remain favorable to the nonlinear ramp. Exact small-kernel calibration confirms schedule-dependent BKS bias at the released threshold; independent cuTensorNet MPS likewise loses the nonlinear BKS advantage when tightened. Strong classical controls dominate in both success rate and runtime. The broader lesson is methodological, not a claim of quantum advantage. When approximate simulation stands between a variational algorithm and a benchmark claim, convergence of the *ranking observable* is part of the algorithm evaluation, not an optional simulator appendix.

References

- [1] E. Farhi, J. Goldstone, and S. Gutmann, “A Quantum Approximate Optimization Algorithm,” arXiv:1411.4028 (2014).
- [2] A. Galda et al., “Similarity-Based Parameter Transferability in the Quantum Approximate Optimization Algorithm,” arXiv:2307.05420 (2023).
- [3] J. A. Montanez-Barrera, D. Willsch, and K. Michielsen, “Transfer learning of optimal QAOA parameters in combinatorial optimization,” Quantum Information Processing 24, 129 (2025).
- [4] S. H. Sureshbabu et al., “Parameter Setting in Quantum Approximate Optimization of Weighted Problems,” Quantum 8, 1231 (2024).
- [5] J. A. Montanez-Barrera et al., “Toward a linear-ramp QAOA protocol: evidence of a scaling advantage in solving some combinatorial optimization problems,” npj Quantum Information 11 (2025).
- [6] F. Ferrari, M. Vandelli, and D. Dragoni, “Feasibility-driven QAOA with penalty scheduling,” arXiv:2606.25117 (2026).
- [7] R. Storn and K. Price, “Differential Evolution—A Simple and Efficient Heuristic for Global Optimization over Continuous Spaces,” Journal of Global Optimization 11, 341–359 (1997).
- [8] T. Koch et al., “The Quantum Optimization Benchmarking Library,” Nature Computational Science 6, 653–671 (2026).
- [9] J. A. Montanez-Barrera, “QAOA Aer/MPS solution for QOBLIB es60fst02,” QOBLIB submission 20260611 (2026), <https://github.com/ZIB-AOPT/QOBLIB>.
- [10] U. Schollwöck, “The density-matrix renormalization group in the age of matrix product states,” Annals of Physics 326, 96–192 (2011).
- [11] R. Watanabe, D. Sels, and J. Tindall, “Tensor network surrogate models for variational quantum computation,” arXiv:2604.20180 (2026).
- [12] A. R. Mazumder, M. Z. Mullath, and H. Tepanyan, “Benchmarking Zero-Setup Quantum Circuit Simulators,” arXiv:2607.09882 (2026).
- [13] G. Kalachev et al., “Pilot-Wave Simulator: Exact Classical Sampling from Ideal and Noisy Quantum Circuits up to Hundreds of Qubits,” Quantum 10, 2173 (2026).
- [14] H. Bayraktar et al., “cuQuantum SDK: A High-Performance Library for Accelerating Quantum Science,” 2023 IEEE International Conference on Quantum Computing and Engineering (QCE), 1050–1061 (2023).